A/B Testing

CMSC320

In this exercise you will experiment with the application of statistical inference in A/B testing. You are a Data Scientist at jsFrameworksRUs and you are tasked with conducting an experiment to measure the effect of a webpage redesign on click rate for a link of interest. You decide to use hypothesis testing to analyze the data you gather from the experiment.

Part 1: Compare to known click rate ($p_A = 0.5$)

In the first case, you assume the click rate for the original version of the page (version A) is $p_A = .5$. The experiment you carry out is pretty simple: show the webpage to $n = 50$ subjects and record whether they click on the link of interest or not. You will use this experiment to estimate your parameter of interest: p_B, the click rate for the new page design (version B).

When you carry out your experiment, you record that $s = 30$ subjects clicked on the link of interest.

Based on our discussion in class, you treat this as $n = 50$ draws from a $\mathrm{Bernoulli}(.5)$ random variable, and use the sample mean $\overline{x} = \frac{1}{n} \sum_{i=1}^{n} x_i = \frac{30}{50} = 0.6$ as your estimate $\hat{p}_B$.

You remember that the hypothesis testing framework is setup in a way where you use your experiment to *reject* the hypothesis that the new design *does not* increase click rate. Therefore, you want to test the (null) hypothesis $p_B \leq p_A = 0.5$ and *reject* it if $p(\overline{X} > \hat{p}_B) \leq \alpha$ under this hypothesis. Remember, α is the rejection level, and we will use $\alpha = 0.05$ here.

To compute $p(\overline{X} > \hat{p}_B)$ under the null hypothesis you will use the normal approximation given by the Central Limit Theorem (CLT).

(a) Derive expressions for $\mathrm{E}\overline{X}$ and $\mathrm{Var}(\overline{X})$ under the null hypothesis in terms of p_A. You will need to use the properties of expectations and variances described below. Here, I give you the derivation for $\mathrm{E}\overline{X}$, you need to do the same for $\mathrm{Var}(\overline{X})$.

$$
\begin{align}
\mathrm{E}\overline{X} &= \mathrm{E}\left[\frac{1}{n}\sum_{i=1}^{n} X_i\right] \tag{1} \\
&= \frac{1}{n}\sum_{i=1}^{n} \mathrm{E}X_i \tag{2} \\
&= \frac{1}{n}(np_A) \tag{3} \\
&= p_A \tag{4}
\end{align}
$$

(b) Based on your derivation, compute values for $\mathrm{E}\overline{X}$ and $\mathrm{Var}(\overline{X})$ based on $p_A = 0.5$ and $n = 50$. Use R to do this.

(c) Using the result above, you can now use the CLT by approximating the distribution of $\overline{X}$ as $N(\mathrm{E}\overline{X}, \sqrt{\mathrm{Var}(\overline{X})})$. Based on this approximation, compute $p(\overline{X} > \hat{p}_B)$. Use the R function `pnorm` to compute this.

(d) Should you reject the null hypothesis $p_B \leq p_A$? Why?

(e) What if you had observed the same $\hat{p}_B = 0.6$ but with $n = 100$ samples. Should you reject the null hypothesis in this case? Why?

(f) What is the *smallest* value $\hat{p}_B$ you would reject the null hypothesis with $n = 100$. Use the `qnorm` function for this. Denote this *smallest* value as q_B.

(g) Based on (f), the smallest detectable improvement for $p_A = 0.5$ with $n = 100$ is then $q_B - p_A$. What is the smallest detectable improvement in your experiment (that is, with $n = 50$)?

Part 2: Compare to known click rate ($p_A = 0.75$)

In this second case, you also assume the click rate for the original version is known, but is $p_A = 0.75$. The data recorded for the experiment is the same. You showed the new design to $n = 50$ subjects and recorded that $s = 30$ clicked on the link of interest.

You want to test the hypothesis $p_B \leq 0.75$ and reject it if $p(\overline{X} > \hat{p}_B) < 0.05$ under this hypothesis. Note the probability in this case is different since $p_A = 0.75$.

(a) What are the values of $\mathrm{E}\overline{X}$ and $\mathrm{Var}(\overline{X})$ under the null hypothesis in this case.

(b) Based on the CLT approximation, compute $p(\overline{X} > \hat{p}_B)$ under the null hypothesis.

(c) Should you reject the null hypothesis $p_B \leq 0.75$? Why?

(d) What if you had observed the same $\hat{p}_B = 0.6$ but with $n = 100$ samples. Should you reject the null hypothesis in this case? Why?

(e) What is the *smallest* value $\hat{p}_B$ you should reject the null hypothesis with $n = 100$. Use the `qnorm` function for this. Denote this *smallest* value as q_B.

(f) Based on (e), the smallest detectable improvement for $p_A = 0.75$ with $n = 100$ is then $q_B - p_A$. What is the smallest detectable improvement in your experiment ($n = 50$)?

Part 3

Consider your answers for parts (1g) and (2f). Is the smallest *detectable* improvement in Question (1g) larger or smaller than in Question (2f)? Explain why this makes sense mathematically.

Part 4: Comparing to estimated click rate p_A.

In this more realistic case you estimate click rates for both page designs in your experiment. The experiment you carry out is as follows: when a customer visits the site, they are randomly (and independently from other customers) shown design A or B, and you record if the click on the link of interest or not. You did this for $n = 100$ customers and recorded the following data:

design	number shown	number clicked
A	$n_A = 55$	$s_A = 35$
B	$n_B = 45$	$s_B = 35$

The null hypothesis we want to test in this case is that $p_B - p_A \leq 0$. That is, that the new design *does not* improve the click rate. How can we use what we know about the CLT in this case?

What we will do is treat estimates using sample means $\hat{p_A} = \overline{X}_A$ and $\hat{p}_B = \overline{X}_B$ as random variables and

define a new random variable $Y = \overline{X}_B - \overline{X}_A$ corresponding to the *difference in click rates* $p_B - p_A$. With that, we derive EY and $\text{Var}(Y)$ under the null hypothesis that $p_B - p_A = 0$.

(a) Derive expressions for EY and $\text{Var}(Y)$ under the null hypothesis in terms of $p_A = p_B = p$. You will need to use the properties of expectations and variances described below. Here, I give you the derivation for EY, you need to do the same for $\text{Var}(Y)$.

$$
\begin{aligned}
EY &= E\left[\overline{X}_B - \overline{X}_A\right] & (5)\\
&= E\overline{X}_B - E\overline{X}_A & (6)\\
&= p_B - p_A & (7)\\
&= 0 & (8)
\end{aligned}
$$

(b) It looks like we will need an estimate of $p_A = p_B = p$ for our CLT approximation. Luckily, under the null hypothesis all $n = 100$ observations from this experiment can be treated as independent identically distributed (iid) draws from a $\text{Bernoulli}(p)$ distribution. Based on this observation, what would be your estimate of $p_A = p_B = p$?

(c) Now that you have an estimate of p, compute a value for $\text{Var}(Y)$.

(d) What is your estimate $\hat{y}$ of $p_B - p_A$ based on the data your recorded for this experiment?

Now, we can reject the null hypothesis of no improvement if $p(Y > \hat{y}) \leq \alpha$ under the null hypothesis.

(e) Using the CLT approximation, what is $p(Y > \hat{y})$

(f) Can you reject the null hypothesis of no improvement in this case? Why? Remember, we are using $\alpha = 0.05$.

Bonus: Smallest detectable improvement for estimated click rates

We could compute smallest detectable improvements in parts 1 and 2 above because we assumed p_A was known. For part 4, we don't know p_A and instead estimate it, so we cannot compute a smallest detectable improvement before the experiment is run because we don't know $p_B = p_A = p$. We can however, compute what the smallest detectable difference *would be* for different values of p.

(a) Make a line plot, with p in the x-axis and the smallest detectable difference as a function of p in the y-axis. You should assume $n_A = 55$ and $n_B = 45$ as above. Again, use the `qnorm` function for this.

Expectation and variance properties

Properties of expectation

(i) $E(aX) = aEX$ for constant a and random variable X
(ii) $E(X + Y) = EX + EY$ for random variables X and Y

Properties of variance

(i) $\text{Var}(aX) = a^2\text{Var}(X)$ for constant a and random variable X
(ii) $Var(X + Y) = Var(X) + Var(Y)$ for *independent* random variables X and Y

Submission

Prepare an Rmarkdown file with your derivations and answer, including code you used to get your answers. Knit to PDF (or save HTML to PDF) and submit to ELMS.